

Sample-Efficient Reinforcement Learning of Partially Observable Markov Games

Qinghua Liu*

Csaba Szepesvári[†]

Chi Jin[‡]

Abstract

This paper considers the challenging tasks of Multi-Agent Reinforcement Learning (MARL) under partial observability, where each agent only sees her own individual observations and actions that reveal incomplete information about the underlying state of system. This paper studies these tasks under the general model of multiplayer general-sum Partially Observable Markov Games (POMGs), which is significantly larger than the standard model of Imperfect Information Extensive-Form Games (IIIEFGs). We identify a rich subclass of POMGs—weakly revealing POMGs—in which sample-efficient learning is tractable. In the self-play setting, we prove that a simple algorithm combining optimism and Maximum Likelihood Estimation (MLE) is sufficient to find approximate Nash equilibria, correlated equilibria, as well as coarse correlated equilibria of weakly revealing POMGs, in a polynomial number of samples when the number of agents is small. In the setting of playing against adversarial opponents, we show that a variant of our optimistic MLE algorithm is capable of achieving sublinear regret when being compared against the optimal maximin policies. To our best knowledge, this work provides the first line of sample-efficient results for learning POMGs.

1 Introduction

This paper studies Multi-Agent Reinforcement Learning (MARL) under *partial observability*, where each player tries to maximize her own utility via interacting with an unknown environment as well as other players. In addition, each agent only sees her own observations and actions, which reveal incomplete information about the underlying state of system. A large number of real-world applications can be cast into this framework: in Poker, cards in a player’s hand are hidden from the other players; in many real-time strategy games, players have only access to their local observations; in multi-agent robotic systems, agents with first-person cameras have to cope with noisy sensors and occlusions. While practical MARL systems have achieved remarkable success in a set of partially observable problems including Poker (Brown and Sandholm, 2019), Starcraft (Vinyals et al., 2019), Dota (Berner et al., 2019) and autonomous driving (Shalev-Shwartz et al., 2016), the theoretical understanding of MARL under partial observability remains very limited.

The combination of partial observability with multiagency introduces a number of unique challenges. The non-Markovian nature of the observations forces the agent to maintain memory and reason about beliefs of the system state, all while exploring to collect information about the environment. Consequently, well-known complexity-theoretic results show that learning and planning in partially observable environments is statistically and computationally intractable even in the single-agent setting (Papadimitriou and Tsitsiklis,

*Princeton University. Email: qinghual@princeton.edu

[†]DeepMind and University of Alberta. Email: szepesva@ualberta.ca

[‡]Princeton University. Email: chij@princeton.edu

1987; Mundhenk et al., 2000; Vlassis et al., 2012; Mossel and Roch, 2005). The presence of interaction between multiple agents further complicates the partially observable problems. In addition to dealing with the adaptive nature of other players who can adjust their strategies according to the learner’s past behaviors, the learner is further required to discover and exploit the information asymmetry due to the separate observations of each agent.

As a result, prior theoretical works on partially observable MARL have been mostly focused on a small subset of problems with strong structural assumptions. For instance, the line of works on Imperfect Information Extensive-Form Games (IIEFG) (see, e.g., Zinkevich et al., 2007; Gordon, 2007; Farina and Sandholm, 2021; Kozuno et al., 2021) assumes tree-structured transition with small depth¹ as well as a special type of emission which can be represented as *information sets*. In contrast, this paper considers a more general mathematical model known as Partially Observable Markov Games (POMGs). POMGs are the natural extensions of both Partially Observable Markov Decision Processes (POMDPs)—the standard model for single-agent partially observable RL, and Markov Games (Shapley, 1953)—the standard model for fully observable MARL. Despite the complexity barriers of learning partially observable systems apply to POMGs, they are of a worst-case nature, which do not preclude efficient algorithms for learning interesting subclasses of POMGs. This motivates us to ask the following question:

Can we develop efficient algorithms that learn a rich class of POMGs?

In this paper, we provide the first positive answer to the highlighted question in terms of the *sample efficiency*.² We identify a rich family of tractable POMGs—*weakly revealing* POMGs (see Section 3). The weakly revealing condition only requires the joint observations of all agents to reveal certain amount of information about the latent states, which is satisfied in many real-world applications. The condition rules out the pathological instances where no player has any information to distinguish latent states, which prevents efficient learning in the worst case.

In the self-play setting where the algorithm can control all the players to learn the equilibria by playing against itself, this paper proposes a new simple algorithm—*Optimistic Maximum Likelihood Estimation for Learning Equilibria* (OMLE-Equilibrium). As the name suggests, it combines optimism, MLE principles with equilibria finding subroutines. The algorithm provably finds approximate Nash equilibria, coarse correlated equilibria and correlated equilibria of any weakly-revealing POMGs using a number of samples polynomial in all relevant parameters.

In the setting of playing against adversarial opponents, we measure the performance of our algorithm by comparing against the optimal maximin policies. We first prove that learning in this setting is hard if each player can only see her own observations and actions. Nevertheless, if the agent is allowed to access other players’ observations and actions *after* each episode of play (e.g., watch the replays of the games from other players’ perspectives afterwards), then we can design a new algorithm OMLE-Adversary which achieves sublinear regret.

To our best knowledge, this is the first line of provably sample-efficient results for learning rich classes of POMGs. Importantly, the classes of problems that can be learned in this paper are significantly larger than known tractable classes of MARL problems under partial observability.

¹The sample complexity of learning IIEFGs scale polynomially with respect to the number of information sets, which typically has an exponential growth in depth.

²For computational efficiency, due to the inherent hardness of planning in POMDPs, all existing provable algorithms that learn large classes of POMDPs (single-agent version of POMGs) require super-polynomial time. We leave the challenge of computationally efficient learning for future work.

1.1 Technical novelty

This paper builds upon the recent progress in learning single-agent POMDPs (Liu et al., 2022a), which identifies the class of weakly revealing POMDPs and develops OMLE algorithm for learning the optimal policy. Besides obtaining a completely new set of results in the multi-agent setting, we here highlight a few contributions and technical novelties of this paper comparing to Liu et al. (2022a).

- This paper rigorously formulates the models, related concepts and learning objectives of multi-player general-sum POMGs, and provides the first line of sample-efficient learning results.
- Extending the weakly revealing conditions into the multi-agent setting lead to two natural candidates: either (a) joint observations or (b) individual observations are required to weakly reveal the state information. This paper shows the former (the weaker assumption) suffices to guarantee tractability.
- Results in the self-play setting requires careful design of optimistic planning algorithms that effectively address the game-theoretical aspects of the problem under partial observability. We achieve this by Subroutine 1, which is even distinct from the standard techniques for learning MGs.
- The discussions and results in the setting of playing against adversarial opponents are completely new, and unique to the multi-agent setup.

1.2 Related works

Reinforcement learning has been extensively studied in the single-agent fully-observable setting (see, e.g., Azar et al., 2017; Dann et al., 2017; Jin et al., 2018, 2020a; Zanette et al., 2020; Jiang et al., 2017; Jin et al., 2021a, and the references therein). For the purpose of this paper, we focus on reviewing existing works on partially observable RL and multi-agent RL in the *exploration* setting.

Markov games In recent years, there has been growing interest in studying Markov games (Shapley, 1953) — the standard generalization of MDPs from the single-player setting to the multi-player setting. Various sample-efficient algorithms have been designed for either two-player zero-sum MGs (e.g., Brafman and Tennenholtz, 2002; Wei et al., 2017; Bai and Jin, 2020; Liu et al., 2021; Bai et al., 2020; Xie et al., 2020; Jin et al., 2021b) or multi-player general-sum MGs (e.g., Liu et al., 2021; Jin et al., 2021c; Song et al., 2021; Daskalakis et al., 2022). However, all these works rely on the states being fully observable, while the POMGs studied in this paper allow states to be only partially observable, which strictly generalizes MGs.

POMDPs POMDPs generalize MDPs from the fully observable setting to the partially observable setting. It is well-known that in POMDPs both planning (Papadimitriou and Tsitsiklis, 1987; Vlassis et al., 2012) and model estimation (Mossel and Roch, 2005) are computationally hard in the worst case. Besides, reinforcement learning of POMDPs is also known to be statistically hard: Krishnamurthy et al. (2016) proved that finding a near-optimal policy of a POMDP in the worst case requires a number of samples that is exponential in the episode length. The hard instances are those pathological POMDPs where the observations contain no useful information for identifying the system dynamics. Nonetheless, these hardness results are all in the worst-case sense and there are still many intriguing positive results on sample-efficient learning of subclasses of POMDPs. For example, Guo et al. (2016); Azizzadenesheli et al. (2016); Jin et al. (2020b); Cai et al. (2022) applied the method of moments to learning undercomplete POMDPs and Liu et al. (2022a) developed the optimistic MLE approach for learning both undercomplete and overcomplete POMDPs. We refer interested readers to Liu et al. (2022a) for a thorough review of existing results on POMDPs.

In terms of algorithmic design, our algorithms build upon the optimistic MLE methodology developed in Liu et al. (2022a). Compared to Liu et al. (2022a), our main algorithmic contribution lies in the design of the optimistic equilibrium computation subroutine in the self-play setting and the optimistic maximin policy design in the adversarial setting. In terms of analysis, our proofs require new techniques tailored to controlling game-theoretic regret, in addition to the OMLE guarantees imported from Liu et al. (2022a). For more detailed explanations of our technical contribution, please refer to Section 1.1.

Imperfect-information extensive-form games In the literature on game theory, there is a long history of learning Imperfect-Information Extensive-Form Games (IIEFGs), (see, e.g., Zinkevich et al., 2007; Gordon, 2007; Farina et al., 2020; Farina and Sandholm, 2021; Kozuno et al., 2021) and the references therein. IIEFGs can be viewed as special cases of POMGs with *tree-structured* transition and a special form of emission which can be represented as information sets. Such emission structure prohibits any state from emitting a distribution over *more than one* observations. Consequently, their results do not apply to learning POMGs with general transition and emission.

Decentralized POMDPs There is another classic model for studying multi-agent partially observable RL, named decentralized POMDPs (e.g., Oliehoek, 2012; Oliehoek and Amato, 2016), which is a special subclass of POMGs where all players share a common reward target. Compared to general POMGs, decentralized POMDPs can only simulate cooperative relations among players, while general POMGs can model both cooperative and competitive relations. Besides, most works (e.g., Nair et al., 2003; Bernstein et al., 2005; Oliehoek et al., 2008; Szer et al., 2012; Dibangoye et al., 2016; Liu et al., 2017; Amato et al., 2019) along this direction mainly focus on the computational complexity of planning with *known* models or simulators instead of the sample efficiency of learning from interactions (as in this paper), which requires learning and estimating the unknown environment while balancing the tradeoff between exploration and exploitation.

Interactive POMDPs Another related model is Interactive POMDPs (I-POMDPs) (Gmytrasiewicz and Doshi, 2005; Doshi et al., 2009), which generalizes POMDPs to handle the presence of other agents by augmenting the latent state with behavior models of other agents. Since I-POMDPs are strictly more general than POMDPs, they are also less understood in theory than POMDPs. Noticably, I-POMDPs can only efficiently handle the problems where the number of possible models of other agents is not too large. For the POMGs model considered in this paper, even in the simplest setting of two-player zero-sum games, an agent can choose from doubly exponentially many different history-dependent strategies. As a result, I-POMDPs cannot efficiently simulate POMGs without using a latent state space that is doubly exponentially large.

2 Preliminary

In this paper, we consider Partially Observable Markov Games (POMGs) in its most generic—multiplayer general-sum form. Formally, we denote a tabular episodic POMG with n players by tuple $(H, \mathcal{S}, \{\mathcal{A}_i\}_{i=1}^n, \{\mathcal{O}_i\}_{i=1}^n; \mathbb{T}, \mathbb{O}, \mu_1; \{r_i\}_{i=1}^n)$, where H denotes the length of each episode, \mathcal{S} the state space with $|\mathcal{S}| = S$, \mathcal{A}_i denotes the action space for the i^{th} player with $|\mathcal{A}_i| = A_i$. We denote by $\mathbf{a} := (a_1, \dots, a_n)$ the joint actions of all n players, and by $\mathcal{A} := \mathcal{A}_1 \times \dots \times \mathcal{A}_n$ the joint action space with $|\mathcal{A}| = A = \prod_i A_i$. $\mathbb{P} = \{\mathbb{P}_h\}_{h \in [H]}$ is the collection of transition matrices, so that $\mathbb{P}_h(\cdot | s, \mathbf{a}) \in \Delta_{\mathcal{S}}$ gives the distribution of the next state if joint actions \mathbf{a} are taken at state s at step h . μ_1 denotes the distribution of the initial state s_1 . \mathcal{O}_i

denotes the observation space for the i^{th} player with $|\mathcal{O}_i| = O_i$. We denote by $\mathbf{o} := (o_1, \dots, o_n)$ the joint observations of all n players, and by $\mathcal{O} := \mathcal{O}_1 \times \dots \times \mathcal{O}_n$ with $|\mathcal{O}| = O = \prod_i O_i$. $\mathbb{O} = \{\mathbb{O}_h\}_{h \in [H]} \subseteq \mathbb{R}^{O \times S}$ is the collection of joint emission matrices, so that $\mathbb{O}_h(\cdot | s) \in \Delta_{\mathcal{O}}$ gives the emission distribution over the joint observation space \mathcal{O} at state s and step h . Finally $r_i = \{r_{i,h}\}_{h \in [H]}$ is the collection of known reward functions for the i^{th} player, so that $r_{i,h}(o_i) \in [0, 1]$ gives the deterministic reward received by the i^{th} player if she observes o_i at step h .³ We remark that since the relation among the rewards of different players can be arbitrary, this model of POMGs subsumes both the cooperative and the competitive settings in partially observable MARL.

In a POMG, the states are always hidden from all players, and each player only observes **her own individual observations and actions**. That is, each player can not see the observations and actions of the other players. At the beginning of each episode, the environment samples s_1 from μ_1 . At each step $h \in [H]$, each player i observes her own observation $o_{i,h}$ where $\mathbf{o}_h := (o_{1,h}, \dots, o_{n,h})$ are jointly sampled from $\mathbb{O}_h(\cdot | s_h)$. Then each player i receives reward $r_{i,h}(o_{i,h})$ and picks action $a_{i,h} \in \mathcal{A}_i$ simultaneously. After that the environment transitions to the next state $s_{h+1} \sim \mathbb{P}_h(\cdot | s_h, \mathbf{a}_h)$ where $\mathbf{a}_h := (a_{1,h}, \dots, a_{n,h})$. The current episode terminates immediately once s_{H+1} is reached.

Policy, value function To define different types of policies, we extend the conventions in *fully observable* Markov games (Jin et al., 2021c) to the partially observable settings. A (*random*) *policy* π_i of the i^{th} player is a map $\pi_i : \Omega \times \bigcup_{h=1}^H ((\mathcal{O}_i \times \mathcal{A}_i)^{h-1} \times \mathcal{O}_i) \rightarrow \mathcal{A}_i$, which maps a random seed ω from space Ω and a history of length $h \in [H]$ —say $\tau_{i,h} := (o_{i,1}, a_{i,1}, \dots, o_{i,h})$, to an action in \mathcal{A}_i . To execute policy π_i , we first draw a random sample ω at the beginning of the episode. Then, at each step h , the i^{th} player simply takes action $\pi_i(\omega, \tau_{i,h})$. We note here ω is shared among all steps $h \in [H]$. ω encodes both the correlation among steps and the individual randomness of each step. We further say a policy π_i is *deterministic* if $\pi_i(\omega, \tau_{i,h}) = \pi_i(\tau_{i,h})$ which is independent of the choice of ω .

By definition, a random policy is equivalent to a mixture of deterministic policies because given a fixed ω the decision of π_i on any history is deterministic. With slight abuse of notation, we use $\pi_i(\omega, \cdot)$ to refer to the deterministic policy realized by policy π_i and a fixed ω . We denote the set of all policies of the i^{th} player by Π_i and the set of all deterministic ones by Π_i^{det} .

A *joint (potentially correlated) policy* is a set of policies $\{\pi_i\}_{i=1}^n$, where the same random seed ω is shared among all agents, which we denote as $\pi = \pi_1 \odot \pi_2 \odot \dots \odot \pi_n$. We also denote $\pi_{-i} = \pi_1 \odot \dots \odot \pi_{i-1} \odot \pi_{i+1} \odot \dots \odot \pi_n$ to be the joint policy excluding the i^{th} player. A special case of joint policy is the *product policy* where the random seed has special form $\omega = (\omega_1, \dots, \omega_n)$, and for any $i \in [n]$, π_i only uses the randomness in ω_i , which is independent of remaining $\{\omega_j\}_{j \neq i}$, which we denote as $\pi = \pi_1 \times \pi_2 \times \dots \times \pi_n$.

We define the value function V_i^π as the expected cumulative reward that the i^{th} player will receive if all players follow joint policy π :

$$V_i^\pi := \mathbb{E}_\pi \left[\sum_{h=1}^H r_{i,h}(o_{i,h}) \right]. \quad (1)$$

where the expectation is taken over the randomness in the initial state, the transitions, the emissions, and the random seed ω in policy π .

Best response and strategy modification For any strategy π_{-i} , the *best response* of the i^{th} player is defined as a policy of the i^{th} player, which is independent of the randomness in π_{-i} and achieves the highest

³This is equivalent to assuming the reward information is contained in the observation.

value for herself conditioned on all other players deploying π_{-i} . Formally, the best response is the maximizer of $\max_{\pi'_i} V_i^{\pi'_i \times \pi_{-i}}$ whose value we also denote as $V_i^{\dagger, \pi_{-i}}$ for simplicity. By its definition, we know the best response can always be achieved by *deterministic* policies.

A *strategy modification* for the i^{th} player is a map $\phi_i : \Pi_i^{\text{det}} \rightarrow \Pi_i^{\text{det}}$, which maps a deterministic policy in Π_i^{det} to another one in it. For any such strategy modification ϕ_i , we can naturally extend its domain and image to include random policies, i.e., define its extension $\phi_i : \Pi_i \rightarrow \Pi_i$ as follows: by definition, a random policy π can be expressed as a mixture of deterministic policies, i.e., as $\pi(\omega, \cdot)$ (a deterministic policy for a fixed ω) with a distribution over ω . Then if we apply map ϕ_i on random policy π , we can define the resulting random policy (denoted as $\phi_i \diamond \pi_i$) as $\phi_i(\pi_i(\omega, \cdot))$ (again a deterministic policy for a fixed ω) with the same distribution over ω . For any joint policy π , we define the best strategy modification of the i^{th} player as the maximizer of $\max_{\phi_i} V_i^{(\phi_i \diamond \pi_i) \odot \pi_{-i}}$.

Different from the best response, which is completely independent of the randomness in π_{-i} , the best strategy modification changes the policy of the i^{th} player while still utilizing the shared randomness among π_i and π_{-i} . Therefore, the best strategy modification is more powerful than the best response: formally one can show that $\max_{\phi_i} V_i^{(\phi_i \diamond \pi_i) \odot \pi_{-i}} \geq \max_{\pi'_i} V_i^{\pi'_i \times \pi_{-i}}$ for any policy π .

2.1 Learning objectives

We focus on three classic equilibrium concepts in game theory—Nash Equilibrium, Correlated Equilibrium (CE) and Coarse Correlated Equilibrium (CCE). First, a Nash equilibrium is defined as a product policy in which no player can increase her value by changing only her own policy. Formally,

Definition 1 (Nash Equilibrium). A *product* policy π is a **Nash equilibrium** if $V_i^{\dagger, \pi_{-i}} = V_i^\pi$ for all $i \in [n]$. A *product* policy π is an ϵ -approximate Nash equilibrium if $V_i^{\dagger, \pi_{-i}} \leq V_i^\pi + \epsilon$ for all $i \in [n]$.

The Nash-regret of a sequence of product policies is the cumulative violation of the Nash condition.

Definition 2 (Nash-regret). Let π^k denote the (product) policy deployed by an algorithm in the k^{th} episode. After a total of K episodes, the Nash-regret is defined as

$$\text{Regret}_{\text{Nash}}(K) = \sum_{k=1}^K \max_{i \in [n]} (V_i^{\dagger, \pi_{-i}^k} - V_i^{\pi^k}).$$

Second, a coarse correlated equilibrium is defined as a joint (potentially correlated) policy where no player can increase her value by unilaterally changing her own policy. Formally,

Definition 3 (Coarse Correlated Equilibrium). A *joint* policy π is a **CCE** if $V_i^{\dagger, \pi_{-i}} \leq V_i^\pi$ for all $i \in [n]$. A *joint* policy π is an ϵ -approximate CCE if $V_i^{\dagger, \pi_{-i}} \leq V_i^\pi + \epsilon$ for all $i \in [n]$.

The only difference between Definition 1 and Definition 3 is that a Nash equilibrium has to be a product policy while a CCE can be correlated. Therefore, CCE is a relaxed notion of Nash equilibrium, and a Nash equilibrium is always a CCE. Similarly, we can define the CCE-regret for a sequence of potentially correlated policies as the cumulative violation of the CCE condition.

Definition 4 (CCE-regret). Let π^k denote the policy deployed by an algorithm in the k^{th} episode. After a total of K episodes, the CCE-regret is defined as

$$\text{Regret}_{\text{CCE}}(K) = \sum_{k=1}^K \max_{i \in [n]} (V_i^{\dagger, \pi_{-i}^k} - V_i^{\pi^k}).$$

Finally, a correlated equilibrium is defined as a joint (potentially correlated) policy where no player can increase her value by unilaterally applying any strategy modification. Formally,

Definition 5 (Correlated Equilibrium). A joint policy π is a **CE** if $\max_{\phi_i} V_i^{(\phi_i \diamond \pi_i) \odot \pi_{-i}} = V_i^\pi$ for all $i \in [n]$. A joint policy π is an ϵ -approximate CE if $\max_{\phi_i} V_i^{(\phi_i \diamond \pi_i) \odot \pi_{-i}} \leq V_i^\pi + \epsilon$ for all $i \in [n]$.

In Partially Observable Markov games, we always have that a Nash equilibrium is a CE, and a CE is a CCE. Finally, we define the CE-regret to be the cumulative violation of the CE condition.

Definition 6 (CE-regret). Let π^k denote the policy deployed by an algorithm in the k^{th} episode. After a total of K episodes, the CE-regret is defined as

$$\text{Regret}_{\text{CE}}(K) = \sum_{k=1}^K \max_{i \in [n]} \max_{\phi_i} (V_i^{(\phi_i \diamond \pi_i^k) \odot \pi_{-i}^k} - V_i^{\pi^k}).$$

3 Weakly Revealing Partially Observable Markov Games

In this section, we define the class of weakly revealing POMGs. To begin with, we consider undercomplete POMGs where there are more observations than hidden states, i.e., $O \geq S$. Formally, the family of α -weakly revealing POMGs includes all POMGs, in which the S^{th} singular value of each emission matrix \mathbb{O}_h is lower bounded by $\alpha > 0$.

Assumption 1 (α -weakly revealing condition). There exists $\alpha > 0$, such that $\min_h \sigma_S(\mathbb{O}_h) \geq \alpha$.

Assumption 1 is simply a robust version of the condition that the rank of each emission matrix is S , which guarantees that no two different latent state mixtures can generate the same observation distribution, i.e., $\mathbb{O}_h \nu_1 \neq \mathbb{O}_h \nu_2$ for any different $\nu_1, \nu_2 \in \Delta_S$. Intuitively, this guarantees that the **joint observations** of all agents contain sufficient information to distinguish any two different state mixtures. We remark that this is much weaker than requiring the individual observations of each agent contain sufficient information about the latent states. The weakly revealing condition is important in excluding those pathological POMGs where the observations contain no useful information for identifying the key parts of model dynamics.

Note that Assumption 1 never holds in the overcomplete setting ($S > O$) as it is impossible to distinguish any two latent state mixtures by only inspecting the observation distribution in a single step. To address this issue, we can instead inspect the observations for m -consecutive steps. To proceed, we define the m -step emission-action matrices

$$\{\mathbb{M}_h \in \mathbb{R}^{(A^{m-1}O^m) \times S}\}_{h \in [H-m+1]}$$

as follows: Given an observation sequence \bar{o} of length m , initial state s and action sequence \bar{a} of length $m-1$, we let $[\mathbb{M}_h]_{(\bar{a}, \bar{o}), s}$ be the probability of receiving \bar{o} provided that the action sequence \bar{a} is used from state s and step h :

$$[\mathbb{M}_h]_{(\bar{a}, \bar{o}), s} = \mathbb{P}(o_{h:h+m-1} = \bar{o} \mid s_h = s, a_{h:h+m-2} = \bar{a}), \quad \forall (\bar{a}, \bar{o}, s) \in \mathcal{A}^{m-1} \times \mathcal{O}^m \times \mathcal{S}. \quad (2)$$

Similar to the undercomplete setting, the weakly-revealing condition in the over-complete setting simply assumes the S^{th} singular value of each m -step emission-action matrix is lower bounded.

Assumption 2 (multistep α -weakly revealing condition). There exists $m \in \mathbb{N}$, $\alpha > 0$ such that $\min_h \sigma_S(\mathbb{M}_h) \geq \alpha$ where \mathbb{M}_h is the m -step emission matrix defined in (2).

Algorithm 1 OMLE-Equilibrium

- 1: **Initialize:** $\mathcal{B}^1 = \{\hat{\theta} \in \Theta : \min_h \sigma_S(\hat{\mathbb{O}}_h) \geq \alpha\}, \mathcal{D} = \{\}$
- 2: **for** $k = 1, \dots, K$ **do**
- 3: compute $\pi^k = \text{Optimistic_Equilibrium}(\mathcal{B}^k)$
- 4: follow π^k to collect a trajectory $\tau^k = (\mathbf{o}_1^k, \mathbf{a}_1^k, \dots, \mathbf{o}_H^k, \mathbf{a}_H^k)$
- 5: add (π^k, τ^k) into \mathcal{D} and update

$$\mathcal{B}^{k+1} = \left\{ \hat{\theta} \in \Theta : \sum_{(\pi, \tau) \in \mathcal{D}} \log \mathbb{P}_{\hat{\theta}}^{\pi}(\tau) \geq \max_{\theta' \in \Theta} \sum_{(\pi, \tau) \in \mathcal{D}} \log \mathbb{P}_{\theta'}^{\pi}(\tau) - \beta \right\} \cap \mathcal{B}^1$$

- 6: output π^{out} that is sampled uniformly at random from $\{\pi^k\}_{k \in [K]}$
-

Subroutine 1 Optimistic_Equilibrium(\mathcal{B})

- 1: **for** $i \in [n]$ **do**
 - 2: let $\bar{V}_i \in \mathbb{R}^{|\Pi_1^{\text{det}}| \times \dots \times |\Pi_n^{\text{det}}|}$ with its π^{th} entry equal to $\sup_{\hat{\theta} \in \mathcal{B}} V_i^{\pi}(\hat{\theta})$ for $\pi \in \Pi_1 \times \dots \times \Pi_n$
 - 3: return EQUILIBRIUM($\bar{V}_1, \dots, \bar{V}_n$)
-

Assumption 2 ensures that m -step consecutive observations shall contain sufficient information to distinguish any two different latent state mixtures. Note that Assumption 1 is a special case of Assumption 2 with $m = 1$. Finally, we remark that the single-agent versions of Assumption 1 and 2 were first identified in Jin et al. (2020b) and Liu et al. (2022a) as sufficient conditions for sample-efficient learning of single-step and multi-step weakly revealing POMDPs (the single-agent version of POMGs), respectively.

4 Learning Equilibria with Self-play

In this section, we study the self-play setting where the algorithm can control all the players to learn the equilibria by playing against itself. We propose a new algorithm — *Optimistic Maximum Likelihood Estimation for Learning Equilibria* (OMLE-Equilibrium) that can provably find Nash equilibria, coarse correlated equilibria and correlate equilibria in any weakly-revealing partially observable Markov games using a number of samples polynomial in all relevant parameters.

4.1 Undercomplete partially observable Markov games

We first present the algorithm and results for learning undercomplete POMGs under Assumption 1. We will see in the later section that with a minor modification the same algorithm also applies to learning overcomplete POMGs under Assumption 2.

Algorithm description To condense notations, we will use $\theta = (\mathbb{T}, \mathbb{O}, \mu_1)$ to denote the parameters of a POMG. Given a policy π and a trajectory τ , we denote by $V_i^{\pi}(\theta)$ the i^{th} player’s value and by $\mathbb{P}_{\theta}^{\pi}(\tau)$ the probability of observing trajectory τ , both under policy π in the POMG model parameterized by θ . We describe OMLE-Equilibrium in Algorithm 1. In each episode, the algorithm executes the following two key steps:

- **Optimistic equilibrium computation** (Line 3) We first invoke Optimistic_Equilibrium

(Subroutine 1) with confidence set \mathcal{B}^k to compute a joint (potentially correlated) policy π^k . Formally, subroutine `Optimistic_Equilibrium`(\mathcal{B}^k) consists of two components:

- **Optimistic value estimation** (Line 1-2 of Subroutine 1) For each player $i \in [n]$ and deterministic joint policy $\pi \in \Pi_1^{\text{det}} \times \dots \times \Pi_n^{\text{det}}$, we compute an upper bound \bar{V}_i^π for the i^{th} player’s value under policy π by using the most optimistic POMG model in the confidence set \mathcal{B}^k .
- **Equilibria computation** (Line 3 of Subroutine 1) Given the optimistic value estimates for all deterministic joint policies and all players, we can view the POMG as a normal-form game where the i^{th} player’s pure strategies consist of all her deterministic policies (i.e., Π_i^{det}) and the payoff she receives under a joint deterministic policy $\pi \in \Pi_1^{\text{det}} \times \dots \times \Pi_n^{\text{det}}$ is equal to the corresponding optimistic value estimate \bar{V}_i^π . Then we compute a `EQUILIBRIUM` π^k for this normal-form game, which is a mixture of all the deterministic joint policies in $\Pi_1^{\text{det}} \times \dots \times \Pi_n^{\text{det}}$.
- **Confidence set update** (Line 4-5) We first follow π^k to collect a trajectory, and then utilize the newly collected data to update the model confidence set via MLE principle.

Here we highlight two algorithmic designs in `OMLE-Equilibrium`: the flexibility of equilibrium computation and the MLE confidence set construction. In the step of equilibrium computation, we can choose `EQUILIBRIUM` to be Nash equilibrium or correlated equilibrium (CE) or coarse correlated equilibrium (CCE) of the normal-form game depending on the target type of equilibrium we aim to learn for the POMG. With regard to the confidence set design, we adopt the idea from Liu et al. (2022a) to include all the POMG models whose likelihood on the historical data is close to the maximum likelihood. This can be viewed as a relaxation of the classic MLE method, with the degree of relaxation controlled by parameter β . One important benefit of this relaxation is that although the groundtruth POMG model is in general not a solution of MLE, its likelihood ratio is rather close to the maximal likelihood. By doing so, we can guarantee the true model is included in the confidence set with high probability. Finally, we remark that Algorithm 1 is computationally inefficient in general due to the steps of optimistic value estimation and equilibrium computation.

Theoretical guarantees Below we present the main theorem for `OMLE-Equilibrium`.

Theorem 7. (Regret of `OMLE-Equilibrium`) *Under Assumption 1, there exists an absolute constant c such that for any $\delta \in (0, 1]$ and $K \in \mathbb{N}$, Algorithm 1 with $\beta = c(H(S^2A + SO)\log(SAOHK) + \log(K/\delta))$ and `EQUILIBRIUM` being one of $\{\text{Nash}, \text{CCE}, \text{CE}\}$ satisfies (respectively) that with probability at least $1 - \delta$,*

$$\text{Regret}_{\{\text{Nash}, \text{CCE}, \text{CE}\}}(k) \leq \text{poly}(S, A, O, H, \alpha^{-1}, \log(K\delta^{-1})) \cdot \sqrt{k} \quad \text{for all } k \in [K].$$

Theorem 7 claims that if all players follow `OMLE-Equilibrium`, then the cumulative $\{\text{Nash}, \text{CCE}, \text{CE}\}$ -regret is upper bounded by $\tilde{\mathcal{O}}(\sqrt{k})$ for any weakly-revealing POMGs that satisfy Assumption 1, where the growth rate w.r.t k is optimal. By the standard online-to-batch conversion, it directly implies the following sample complexity result:

Corollary 8. (Sample Complexity of `OMLE-Equilibrium`) *Under the same setting as Theorem 7, when $K \geq \text{poly}(S, A, O, H, \alpha^{-1}, \log(\epsilon^{-1}\delta^{-1})) \cdot \epsilon^{-2}$, then with probability at least $1/2$, π^{out} is an ϵ - $\{\text{Nash}, \text{CCE}, \text{CE}\}$ policy.*

Here the dependence on the precision parameter ϵ is optimal. Finally, notice that the upper bound in Theorem 7 depends polynomially on the inverse of α — a lower bound for the minimal singular value of the joint emission matrix \mathbb{O}_h in Assumption 1. This dependence is shown to be unavoidable even in the single-player setting (POMDPs) (Liu et al., 2022a).

Algorithm 2 multi-step OMLE-Equilibrium

- 1: **Initialize:** $\mathcal{B}^1 = \{\hat{\theta} \in \Theta : \min_h \sigma_S(\hat{\mathbb{M}}_h) \geq \alpha\}$, $\mathcal{D} = \{\}$
 - 2: **for** $k = 1, \dots, K$ **do**
 - 3: compute $\pi^k = \text{Optimistic_Equilibrium}(\mathcal{B}^k)$
 - 4: **for** $h = 0, \dots, H - m$ **do**
 - 5: execute policy $\pi_{1:h}^k \circ \text{uniform}(\mathcal{A})$ to collect a trajectory $\tau^{k,h}$
 then add $(\pi_{1:h}^k \circ \text{uniform}(\mathcal{A}), \tau^{k,h})$ into \mathcal{D}
 - 6: update $\mathcal{B}^{k+1} = \left\{ \hat{\theta} \in \Theta : \sum_{(\pi, \tau) \in \mathcal{D}} \log \mathbb{P}_{\hat{\theta}}^{\pi}(\tau) \geq \max_{\theta' \in \Theta} \sum_{(\pi, \tau) \in \mathcal{D}} \log \mathbb{P}_{\theta'}^{\pi}(\tau) - \beta \right\} \cap \mathcal{B}^1$
 - 7: output π^{out} that is sampled uniformly at random from $\{\pi^k\}_{k \in [K]}$
-

4.2 Overcomplete partially observable Markov games

In this subsection, we extend OMLE-Equilibrium to the more challenging setting of learning overcomplete POMGs, where there can be much less observations than latent states. We prove that a simple variant of OMLE-Equilibrium still enjoys polynomial sample-efficiency guarantee for learning any multi-step weakly revealing POMGs.

Algorithm description We describe the multi-step generalization of OMLE-Equilibrium in Algorithm 2, which inherits the key designs from Algorithm 1 and additionally makes two important modifications to address the challenge of insufficient information from single-step observation. The first change is to utilize a more active sampling strategy for exploration. Instead of simply following the optimistic policy π^k , we will iteratively execute $H - m + 1$ policies of form $\pi_{1:h}^k \circ \text{uniform}(\mathcal{A})$ where the players first follow policy π^k from step 1 to step h , then pick actions uniformly at random to finish the remaining $H - h$ steps. Intuitively, by actively trying random action sequences after executing policy π^k , the algorithm can acquire more information about the system dynamics corresponding to those latent states that are frequently visited by π^k , and therefore help address the challenge of lacking sufficient information from single-step observation. The second change made by Algorithm 2 is that in constructing the confidence set, we require the minimal singular value of the multistep emission matrix to be lower bounded, which enforces the multistep weakly revealing condition in Assumption 2.

Theoretical guarantee Below we present the main theorem for multi-step OMLE-Equilibrium.

Theorem 9. (Total suboptimality of multi-step OMLE-Equilibrium) *Under Assumption 2, there exists an absolute constant c such that for any $\delta \in (0, 1]$ and $K \in \mathbb{N}$, Algorithm 2 with*

$$\beta = c \left(H(S^2 A + SO) \log(SAOHK) + \log(K/\delta) \right)$$

and EQUILIBRIUM being one of $\{\text{Nash}, \text{CCE}, \text{CE}\}$ satisfies (respectively) that with probability at least $1 - \delta$,

$$\text{Regret}_{\{\text{Nash}, \text{CCE}, \text{CE}\}}(k) \leq \text{poly}(S, A^m, O, H, \alpha^{-1}, \log(K\delta^{-1})) \cdot \sqrt{k} \quad \text{for all } k \in [K],$$

where the regret is computed for policy π^1, \dots, π^k .

Theorem 9 claims that the total $\{\text{Nash}, \text{CCE}, \text{CE}\}$ -“regret” (that are computed on policy π^1, \dots, π^k) of multi-step OMLE-Equilibrium is upper bounded by $\tilde{O}(\sqrt{k})$ for any multi-step weakly revealing POMGs

satisfying Assumption 2. We remark that, strictly speaking, Theorem 9 is not a standard regret guarantee since the policies executed by multi-step OMLE-Equilibrium are compositions of π^1, \dots, π^k and random actions, instead of purely π^1, \dots, π^k . Nevertheless, we can still utilize the standard online-to-batch conversion to obtain the following sample complexity guarantee:

Corollary 10. (sample complexity of multi-step OMLE-Equilibrium) *Under the same setting as Theorem 9, when $K \geq \text{poly}(S, A^m, O, H, \alpha^{-1}, \log(\epsilon^{-1}\delta^{-1})) \cdot \epsilon^{-2}$, then with probability at least $1/2$, π^{out} is an ϵ - $\{\text{Nash}, \text{CCE}, \text{CE}\}$ policy.*

Here the dependence on the precision parameter ϵ is optimal up to poly-logarithmic factors. Finally, observe that the sample complexity has A^m dependency that is exponential in m , which is unavoidable in general even in the single-player setting (POMDPs) (Liu et al., 2022a). Nonetheless, in order to make $\min_h \text{rank}(\mathbb{M}_h) = S$ possible (Assumption 2), we only need to make $(OA)^m \gtrsim S$, i.e., $m \gtrsim \log S$ which is very small. In this paper, when we claim the sample complexity is polynomial, we consider m to be small enough so that $A^m \leq \text{poly}(S, A, O, H, \alpha^{-1})$.

5 Playing against Adversarial Opponents

In this section, we turn to the online setting where the learner only controls a single player and the remaining players can execute arbitrary strategies. In this setting, we no longer target at learning game-theoretic equilibria because if other players keep playing some highly suboptimal policies then the learner may never be able to explore the environment thoroughly and thus lacks sufficient information to compute equilibria. Instead, we consider the standard goal for online setting, which is to achieve low regret in terms of cumulative rewards even if all other players play adversarially against the learner. Without loss of generality, we assume the learner only controls the 1st player throughout this section.

5.1 Statistical hardness for the standard setting

We first consider the standard POMG setting where each player can only observe her *own* observations and actions. We prove that achieving low regret in this setting is impossible in general even if (i) the POMG is two-player zero-sum and satisfies Assumption 1 with $\alpha = 1$, (ii) the opponent keeps playing a fixed deterministic policy *known* to the learner, and (iii) the only parts of the model unknown to the learner are the emission matrices.

Theorem 11. *For any $L, k \in \mathbb{N}^+$, there exist (i) a two-player zero-sum POMG of size $S, A, O, H = \mathcal{O}(L)$ and satisfying Assumption 1 with $\alpha = 1$, and (ii) a fixed opponent who keeps playing a known deterministic policy π_2 , so that with probability at least $1/2$*

$$\sum_{t=1}^k \left(\max_{\tilde{\pi}_1} \min_{\tilde{\pi}_2} V_1^{\tilde{\pi}_1 \times \tilde{\pi}_2} - V_1^{\pi_1^t \times \pi_2} \right) \geq \Omega(\min\{2^L, k\}),$$

where π_1^t is the policy played by the learner in the t^{th} episode.

Theorem 11 claims that when the learner is not allowed to access the opponent's observations and actions, there exists exponential regret lower bound for competing with the max-min value (i.e., Nash value in two-player zero-sum POMGs) even in the very benign scenario as described above. We remark that this lower bound directly implies competing with the best fixed policy in hindsight is also hard because the max-min value is always no larger than the value of the best-response to π_2 :

$$\max_{\tilde{\pi}_1} \min_{\tilde{\pi}_2} V_1^{\tilde{\pi}_1 \times \tilde{\pi}_2} \leq \max_{\tilde{\pi}_1} V_1^{\tilde{\pi}_1 \times \pi_2} = V_1^{\dagger, \pi_2}.$$

Algorithm 3 OMLE-Adversary

- 1: **Initialize:** $\mathcal{B}^1 = \{\hat{\theta} \in \Theta : \sigma_S(\hat{\Theta}) \geq \alpha\}$, $\mathcal{D} = \{\}$
- 2: **for** $k = 1, \dots, K$ **do**
- 3: learner computes $(\cdot, \pi_1^k) = \operatorname{argmax}_{\hat{\theta} \in \mathcal{B}^k, \hat{\pi}_1 \in \Pi_1} \min_{\hat{\pi}_{-1} \in \Pi_{-1}} V_1^{\hat{\pi}_1 \times \hat{\pi}_{-1}}(\hat{\theta})$
- 4: opponents pick policies π_{-1}^k
- 5: execute policy $\pi^k = \pi_1^k \times \pi_{-1}^k$ to collect $\tau^k = (\mathbf{o}_1^k, \mathbf{a}_1^k, \dots, \mathbf{o}_H^k, \mathbf{a}_H^k)$
- 6: add (π^k, τ^k) into \mathcal{D} and update

$$\mathcal{B}^{k+1} = \left\{ \hat{\theta} \in \Theta : \sum_{(\pi, \tau) \in \mathcal{D}} \log \mathbb{P}_{\hat{\theta}}^{\pi}(\tau) \geq \max_{\theta' \in \Theta} \sum_{(\pi, \tau) \in \mathcal{D}} \log \mathbb{P}_{\theta'}^{\pi}(\tau) - \beta \right\} \cap \mathcal{B}^1 \quad (3)$$

5.2 Positive results for the game-replay setting

In this section, we consider the game-replay setting where *after* each episode of play, every player will reveal their observations and actions in this episode to other players. In other words, every player is able to observe the whole trajectory $\tau^k = (\mathbf{o}_1^k, \mathbf{a}_1^k, \dots, \mathbf{o}_H^k, \mathbf{a}_H^k)$ *after* the k^{th} episode is finished. The motivation for considering this setting is in many real-world games, e.g., Dota, StarCraft and Poker, players are usually allowed to watch the replays of the games they have played, in which they can freely view other players' observations and actions. Below, we show that a simple variant of OMLE-Equilibrium enjoys sublinear regret when playing against adversarial opponents.

Algorithm description We provide the formal description of OMLE-Adversary in Algorithm 3. Same as OMLE-Equilibrium, OMLE-Adversary utilizes the relaxed MLE approach to construct the confidence set. The key modification lies in the computation of player 1's (stochastic) policy π_1^k (Line 3). Specifically, the learner will compute the most optimistic model θ^k in the confidence set \mathcal{B}^k by examining player 1's max-min value in each model. Then choose player 1's policy to be the one with the highest value under θ^k , assuming all other players jointly play against player 1.

Finally, we remark that although the confidence set construction in Line 6 seems to involve the joint policies π of all players, the confidence set itself is in fact independent of the joint policies π . Therefore, the 1st player can still construct the confidence set without knowledge of other players' policies. This is because the dependency of the loglikelihood function on policy π are equal on both sides of (3), and thus they cancel with each other. Formally, for any $\hat{\theta}, \theta' \in \Theta$, we have

$$\sum_{t=1}^k \left(\log \mathbb{P}_{\hat{\theta}}^{\pi^t}(\tau^t) - \log \mathbb{P}_{\theta'}^{\pi^t}(\tau^t) \right) = \sum_{t=1}^k \left(\log \mathbb{P}_{\hat{\theta}}(\mathbf{o}_{1:H}^t \mid \mathbf{a}_{1:H}^t) - \log \mathbb{P}_{\theta'}(\mathbf{o}_{1:H}^t \mid \mathbf{a}_{1:H}^t) \right).$$

Theoretical guarantees Below we present the main theorem for OMLE-Adversary.

Theorem 12. (Regret of OMLE-Adversary) *Under Assumption 1, there exists an absolute constant c such that for any $\delta \in (0, 1]$ and $K \in \mathbb{N}$, Algorithm 3 with $\beta = c(H(S^2 A + SO) \log(SAOHK) + \log(K/\delta))$ satisfies that with probability at least $1 - \delta$,*

$$\sum_{t=1}^k \left(\max_{\hat{\pi}_1} \min_{\hat{\pi}_2} V_1^{\hat{\pi}_1 \times \hat{\pi}_2} - V_1^{\pi^t} \right) \leq \text{poly}(S, A, O, H, \alpha^{-1}, \log(K\delta^{-1})) \cdot \sqrt{k} \quad \text{for all } k \in [K].$$

Theorem 12 claims that the regret of OMLE-Adversary is upper bounded by $\tilde{O}(\sqrt{k})$ in any weakly revealing POMGs that satisfy Assumption 1, no matter what adversarial strategies other players might take.

Here the regret is defined by comparing the cumulative rewards received by player 1 to the max-min value $\max_{\tilde{\pi}_1} \min_{\tilde{\pi}_2} V_1^{\tilde{\pi}_1 \times \tilde{\pi}_2}$ that is the largest value she could receive if all other players jointly play against her. Notice that this regret is weaker than the typical version of regret considered in online learning literature, which typically competes with the best response in hindsight, i.e.,

$$\max_{\tilde{\pi}_1} \sum_{t=1}^k \left(V_1^{\tilde{\pi}_1 \times \pi_{-1}^t} - V_1^{\pi^t} \right).$$

Therefore, it is natural to ask whether we can obtain similar sublinear regret in terms of the above regret definition. Unfortunately, previous work (Liu et al., 2022b) proved that there exists exponential regret lower bound for competing with the best response in hindsight even in *fully observable* two-player zero-sum Markov games, which are special cases of POMGs satisfying Assumption 1 with $\alpha = 1$. As a result, achieving low regret in the above sense is also intractable in POMGs.

Negative results for generalization to multi-step weakly revealing POMGs So far, we only derive the positive result (Theorem 12) for single-step weakly revealing POMGs. A reader might wonder whether similar results can be obtained in the more general setting of multi-step weakly revealing POMGs. Unfortunately, this generalization turns out to be impossible in general, even if (i) the POMG satisfies Assumption 2 with $m = 2$ and $\alpha \geq 1$, and (ii) the learner can directly observe the opponents’ actions and observations. We defer the formal statement of this hardness result and its proof to Appendix C.3.

References

- Noam Brown and Tuomas Sandholm. Superhuman AI for multiplayer poker. *Science*, 365(6456):885–890, 2019.
- Oriol Vinyals, Igor Babuschkin, Wojciech M Czarnecki, Michael Mathieu, Andrew Dudzik, Junyoung Chung, David H Choi, Richard Powell, Timo Ewalds, Petko Georgiev, et al. Grandmaster level in StarCraft II using multi-agent reinforcement learning. *Nature*, 575(7782):350–354, 2019.
- Christopher Berner, Greg Brockman, Brooke Chan, Vicki Cheung, Przemyslaw Dkebiak, Christy Dennison, David Farhi, Quirin Fischer, Shariq Hashme, Chris Hesse, et al. Dota 2 with large scale deep reinforcement learning. *arXiv preprint arXiv:1912.06680*, 2019.
- Shai Shalev-Shwartz, Shaked Shammah, and Amnon Shashua. Safe, multi-agent, reinforcement learning for autonomous driving. *arXiv preprint arXiv:1610.03295*, 2016.
- Christos H Papadimitriou and John N Tsitsiklis. The complexity of Markov decision processes. *Mathematics of operations research*, 12(3):441–450, 1987.
- Martin Mundhenk, Judy Goldsmith, Christopher Lusena, and Eric Allender. Complexity of finite-horizon Markov decision process problems. *Journal of the ACM (JACM)*, 47(4):681–720, 2000.
- Nikos Vlassis, Michael L Littman, and David Barber. On the computational complexity of stochastic controller optimization in POMDPs. *ACM Transactions on Computation Theory (TOCT)*, 4(4):1–8, 2012.
- Elchanan Mossel and Sébastien Roch. Learning nonsingular phylogenies and hidden Markov models. In *Proceedings of the thirty-seventh annual ACM symposium on Theory of computing*, pages 366–375, 2005.

- Martin Zinkevich, Michael Johanson, Michael Bowling, and Carmelo Piccione. Regret minimization in games with incomplete information. *Advances in neural information processing systems*, 20:1729–1736, 2007.
- Geoffrey J Gordon. No-regret algorithms for online convex programs. In *Advances in Neural Information Processing Systems*, pages 489–496. Citeseer, 2007.
- Gabriele Farina and Tuomas Sandholm. Model-free online learning in unknown sequential decision making problems and games. *arXiv preprint arXiv:2103.04539*, 2021.
- Tadashi Kozuno, Pierre Ménard, Rémi Munos, and Michal Valko. Model-free learning for two-player zero-sum partially observable Markov games with perfect recall. *arXiv preprint arXiv:2106.06279*, 2021.
- Lloyd S Shapley. Stochastic games. *Proceedings of the national academy of sciences*, 39(10):1095–1100, 1953.
- Qinghua Liu, Alan Chung, Csaba Szepesvari, and Chi Jin. When is partially observable reinforcement learning not scary? *arXiv preprint*, 2022a.
- Mohammad Gheshlaghi Azar, Ian Osband, and Rémi Munos. Minimax regret bounds for reinforcement learning. In *International Conference on Machine Learning*, pages 263–272. PMLR, 2017.
- Christoph Dann, Tor Lattimore, and Emma Brunskill. Unifying PAC and regret: Uniform PAC bounds for episodic reinforcement learning. *Advances in Neural Information Processing Systems*, 30, 2017.
- Chi Jin, Zeyuan Allen-Zhu, Sebastien Bubeck, and Michael I Jordan. Is Q-learning provably efficient? *Advances in neural information processing systems*, 31, 2018.
- Chi Jin, Zhuoran Yang, Zhaoran Wang, and Michael I Jordan. Provably efficient reinforcement learning with linear function approximation. In *Conference on Learning Theory*, pages 2137–2143. PMLR, 2020a.
- Andrea Zanette, Alessandro Lazaric, Mykel Kochenderfer, and Emma Brunskill. Learning near optimal policies with low inherent Bellman error. In *International Conference on Machine Learning*, pages 10978–10989. PMLR, 2020.
- Nan Jiang, Akshay Krishnamurthy, Alekh Agarwal, John Langford, and Robert E Schapire. Contextual decision processes with low Bellman rank are PAC-learnable. In *International Conference on Machine Learning*, pages 1704–1713. PMLR, 2017.
- Chi Jin, Qinghua Liu, and Sobhan Miryoosefi. Bellman eluder dimension: New rich classes of RL problems, and sample-efficient algorithms. *Advances in Neural Information Processing Systems*, 34, 2021a.
- Ronen I Brafman and Moshe Tennenholtz. R-max-a general polynomial time algorithm for near-optimal reinforcement learning. *Journal of Machine Learning Research*, 3(Oct):213–231, 2002.
- Chen-Yu Wei, Yi-Te Hong, and Chi-Jen Lu. Online reinforcement learning in stochastic games. In *Advances in Neural Information Processing Systems*, pages 4987–4997, 2017.
- Yu Bai and Chi Jin. Provable self-play algorithms for competitive reinforcement learning. *arXiv preprint arXiv:2002.04017*, 2020.

- Qinghua Liu, Tiancheng Yu, Yu Bai, and Chi Jin. A sharp analysis of model-based reinforcement learning with self-play. In *International Conference on Machine Learning*, pages 7001–7010. PMLR, 2021.
- Yu Bai, Chi Jin, and Tiancheng Yu. Near-optimal reinforcement learning with self-play. In *Advances in Neural Information Processing Systems*, 2020.
- Qiaomin Xie, Yudong Chen, Zhaoran Wang, and Zhuoran Yang. Learning zero-sum simultaneous-move markov games using function approximation and correlated equilibrium. *arXiv preprint arXiv:2002.07066*, 2020.
- Chi Jin, Qinghua Liu, and Tiancheng Yu. The power of exploiter: Provable multi-agent rl in large state spaces. *arXiv preprint arXiv:2106.03352*, 2021b.
- Chi Jin, Qinghua Liu, Yuanhao Wang, and Tiancheng Yu. V-learning—a simple, efficient, decentralized algorithm for multiagent RL. *arXiv preprint arXiv:2110.14555*, 2021c.
- Ziang Song, Song Mei, and Yu Bai. When can we learn general-sum markov games with a large number of players sample-efficiently? *arXiv preprint arXiv:2110.04184*, 2021.
- Constantinos Daskalakis, Noah Golowich, and Kaiqing Zhang. The complexity of markov equilibrium in stochastic games. *arXiv preprint arXiv:2204.03991*, 2022.
- Akshay Krishnamurthy, Alekh Agarwal, and John Langford. PAC reinforcement learning with rich observations. *Advances in Neural Information Processing Systems*, 29, 2016.
- Zhaohan Daniel Guo, Shayan Doroudi, and Emma Brunskill. A PAC RL algorithm for episodic POMDPs. In *Artificial Intelligence and Statistics*, pages 510–518. PMLR, 2016.
- Kamyar Azizzadenesheli, Alessandro Lazaric, and Animashree Anandkumar. Reinforcement learning of POMDPs using spectral methods. In *Conference on Learning Theory*, pages 193–256. PMLR, 2016.
- Chi Jin, Sham M Kakade, Akshay Krishnamurthy, and Qinghua Liu. Sample-efficient reinforcement learning of undercomplete POMDPs. *NeurIPS*, 2020b.
- Qi Cai, Zhuoran Yang, and Zhaoran Wang. Sample-efficient reinforcement learning for pomdps with linear function approximations. *arXiv preprint arXiv:2204.09787*, 2022.
- Gabriele Farina, Christian Kroer, and Tuomas Sandholm. Faster game solving via predictive Blackwell approachability: Connecting regret matching and mirror descent. *arXiv preprint arXiv:2007.14358*, 2020.
- Frans A Oliehoek. Decentralized pomdps. In *Reinforcement Learning*, pages 471–503. Springer, 2012.
- Frans A Oliehoek and Christopher Amato. *A concise introduction to decentralized POMDPs*. Springer, 2016.
- Ranjit Nair, Milind Tambe, Makoto Yokoo, David Pynadath, and Stacy Marsella. Taming decentralized pomdps: Towards efficient policy computation for multiagent settings. In *IJCAI*, volume 3, pages 705–711. Citeseer, 2003.
- Daniel S Bernstein, Eric A Hansen, and Shlomo Zilberstein. Bounded policy iteration for decentralized pomdps. In *Proceedings of the nineteenth international joint conference on artificial intelligence (IJCAI)*, pages 52–57, 2005.

- Frans A Oliehoek, Matthijs TJ Spaan, and Nikos Vlassis. Optimal and approximate q-value functions for decentralized pomdps. *Journal of Artificial Intelligence Research*, 32:289–353, 2008.
- Daniel Szer, François Charpillet, and Shlomo Zilberstein. Maa*: A heuristic search algorithm for solving decentralized pomdps. *arXiv preprint arXiv:1207.1359*, 2012.
- Jilles Steeve Dibangoye, Christopher Amato, Olivier Buffet, and François Charpillet. Optimally solving dec-pomdps as continuous-state mdps. *Journal of Artificial Intelligence Research*, 55:443–497, 2016.
- Miao Liu, Kavinayan Sivakumar, Shayegan Omidshafiei, Christopher Amato, and Jonathan P How. Learning for multi-robot cooperation in partially observable stochastic environments with macro-actions. In *2017 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 1853–1860. IEEE, 2017.
- Christopher Amato, George Konidaris, Leslie P Kaelbling, and Jonathan P How. Modeling and planning with macro-actions in decentralized pomdps. *Journal of Artificial Intelligence Research*, 64:817–859, 2019.
- Piotr J Gmytrasiewicz and Prashant Doshi. A framework for sequential planning in multi-agent settings. *Journal of Artificial Intelligence Research*, 24:49–79, 2005.
- Prashant Doshi, Yifeng Zeng, and Qiongyu Chen. Graphical models for interactive pomdps: representations and solutions. *Autonomous Agents and Multi-Agent Systems*, 18(3):376–416, 2009.
- Qinghua Liu, Yuanhao Wang, and Chi Jin. Learning markov games with adversarial opponents: Efficient algorithms and fundamental limits. *arXiv preprint arXiv:2203.06803*, 2022b.

A Notations

We first introduce some notations that will be frequently used in the remainder of appendix.

- We will use $\mu \in \Pi^{\text{det}}$ to refer to a *deterministic* joint policy, and use $\mu_i \in \Pi_i^{\text{det}}$ to refer to a *deterministic* policy of player i .
- Since each stochastic joint policy $\pi \in \Pi$ is equivalent to a distribution over all the deterministic joint policies Π^{det} , with slight abuse of notation, we denote by $\mu \sim \pi$ the process of sampling a deterministic joint policy μ from the policy distribution specified by π . We can similarly define $\mu_i \sim \pi_i$ for any stochastic policy $\pi_{=i}$ of player i .
- Given a policy π and a POMG model θ , denote by \mathbb{P}_θ^π the distribution over trajectories (i.e., τ_H) produced by executing policy π in a POMG parameterized by θ . Since the reward per trajectory is bounded by H , we always have

$$V_i^\pi(\theta) - V_i^\pi(\hat{\theta}) \leq H \|\mathbb{P}_\theta^\pi - \mathbb{P}_{\hat{\theta}}^\pi\|$$

for any policy π , POMG models $\theta, \hat{\theta}$, and player i .

- Denote by θ^* the parameters of the groundtruth POMG model we are interacting with.

B Proofs for the Self-play Setting

B.1 Proof of Theorem 7

In this section, we prove Theorem 7 with a specific polynomial dependency as stated in the following theorem.

Theorem 13. (Regret of OMLE-Equilibrium) *Under Assumption 1, there exists an absolute constant c such that for any $\delta \in (0, 1]$ and $K \in \mathbb{N}$, Algorithm 1 with $\beta = c(H(S^2A + SO) \log(SAOHK) + \log(K/\delta))$ and EQUILIBRIUM being one of $\{Nash, CCE, CE\}$ satisfies (respectively) that with probability at least $1 - \delta$,*

$$\text{Regret}_{\{Nash, CCE, CE\}}(k) \leq \tilde{O}\left(\frac{S^2AO}{\alpha^2} \sqrt{k(S^2A + SO)} \times \text{poly}(H)\right) \quad \text{for all } k \in [K].$$

The proof consists of three steps:

1. First we rewrite Algorithm 1 in an equivalent form that is perfectly compatible with the analysis in Liu et al. (2022a).
2. After that we can directly import the theoretical guarantees from Liu et al. (2022a) and obtain a sublinear upper bound for the cumulative error of density estimation.
3. Finally, we combine the game-theoretic analysis tailored for POMGs with the density estimation guarantee derived in the second step, which gives the desired sublinear game-theoretic regret.

Algorithm 4 OMLE-Equilibrium

- 1: **Initialize:** $\mathcal{B}^1 = \{\hat{\theta} \in \Theta : \min_h \sigma_S(\hat{\mathbb{O}}_h) \geq \alpha\}, \mathcal{D} = \{\}$
- 2: **for** $k = 1, \dots, K$ **do**
- 3: compute $\pi^k = \text{Optimistic_Equilibrium}(\mathcal{B}^k)$
- 4: sample a deterministic joint policy μ^k from π^k , then follow μ^k to collect a trajectory τ^k
- 5: add (μ^k, τ^k) into \mathcal{D} and update

$$\mathcal{B}^{k+1} = \left\{ \hat{\theta} \in \Theta : \sum_{(\pi, \tau) \in \mathcal{D}} \log \mathbb{P}_{\hat{\theta}}^{\pi}(\tau) \geq \max_{\theta' \in \Theta} \sum_{(\pi, \tau) \in \mathcal{D}} \log \mathbb{P}_{\theta'}^{\pi}(\tau) - \beta \right\} \cap \mathcal{B}^1$$

B.1.1 Step 1

To begin with, we make the following observations about Algorithm 1:

- The sampling procedure in each episode k is equivalent to: first sample a *deterministic* joint policy μ^k from π^k and then execute μ^k to collect a trajectory τ^k .
- In constructing the confidence set \mathcal{B}^k , we can replace π^k with μ^k without making any difference, because the dependency of the loglikelihood function on policy π are equal on both sides of the inequality in \mathcal{B}^k and thus they cancel with each other. Formally, for any $\hat{\theta}, \theta' \in \Theta$, we have

$$\begin{aligned} & \sum_{t=1}^k \left(\log \mathbb{P}_{\hat{\theta}}^{\pi^t}(\tau^t) - \log \mathbb{P}_{\theta'}^{\pi^t}(\tau^t) \right) \\ &= \sum_{t=1}^k \left(\log \mathbb{P}_{\hat{\theta}}(\mathbf{o}_{1:H}^t \mid \mathbf{a}_{1:H}^t) - \log \mathbb{P}_{\theta'}(\mathbf{o}_{1:H}^t \mid \mathbf{a}_{1:H}^t) \right) = \sum_{t=1}^k \left(\log \mathbb{P}_{\hat{\theta}}^{\mu^t}(\tau^t) - \log \mathbb{P}_{\theta'}^{\mu^t}(\tau^t) \right). \end{aligned}$$

Based on the above two observations, Algorithm 1 can be *equivalently* written in the form of Algorithm 4 where we highlight the modified parts in blue.

Remark 14. *The technical reason for rewriting Algorithm 1 in the form of Algorithm 4 is that in the optimistic equilibrium subroutine (Subroutine 1) we utilize the optimistic value estimate for each deterministic joint policy to construct the optimistic normal-form game and compute the optimistic game-theoretic equilibria. As a result, in order to control the cumulative regret due to over-optimism, we need guarantees on the accuracy of optimistic value estimates for deterministic joint policies. This is why we want to explicitly insert the “dummy” deterministic policy μ^k in each episode.*

B.1.2 Step 2

Now we can directly instantiate the analysis of optimistic MLE (Appendix E in Liu et al. (2022a)) on Algorithm 4, which gives the following theoretical guarantee:

Theorem 15. (Liu et al. (2022a)) *Under Assumption 1 and the same choice of β as in Theorem 7, with probability at least $1 - \delta$, Algorithm 4 satisfies that for **all** $k \in [K]$ and **all** $\theta^1 \in \mathcal{B}^1, \dots, \theta^K \in \mathcal{B}^K$*

- $\theta^* \in \mathcal{B}^k$,

- $\sum_{t=1}^k \|\mathbb{P}_{\theta^t}^{\mu^t} - \mathbb{P}_{\theta^*}^{\mu^t}\|_1 \leq \tilde{O}\left(\frac{S^2 AO}{\alpha^2} \sqrt{k(S^2 A + SO)} \times \text{poly}(H)\right).$

We comment that there are two differences between the optimistic MLE algorithm in Liu et al. (2022a) and the Algorithm 4 here: (i) the former one is designed for single-player POMGs, i.e., POMDPs while the latter one is for multi-player POMGs; (ii) μ^t is computed using different criteria. Nonetheless, we can still reuse their theoretical guarantees proved in their Appendix E without making any change because: (i) in the self-play setting, multi-player POMGs can be viewed as POMDPs with a single meta-player whose action space is of cardinality $A = A_1 \times \dots \times A_n$ and observation space is of cardinality $O = O_1 \times \dots \times O_n$; (ii) when proving the second statement in Theorem 15, Liu et al. (2022a) only use the fact that τ^t is sampled from μ^t but allow both μ^t and $\theta^t \in \mathcal{B}^t$ to be arbitrarily chosen. (The only place Liu et al. (2022a) need to use how μ^t and θ^t is computed is in relating the regret to $\sum_{t=1}^k \|\mathbb{P}_{\theta^t}^{\mu^t} - \mathbb{P}_{\theta^*}^{\mu^t}\|_1$, which has nothing to do with the proof of Theorem 15.)

B.1.3 Step 3

Now let us prove Theorem 13 conditioning on the two relations stated in Theorem 15 being true. To proceed, we define

$$\bar{V}_i^{k,\mu} = \max_{\hat{\theta} \in \mathcal{B}^k} V_i^\mu(\hat{\theta}) \quad \text{for any } (\mu, k, i) \in \Pi^{\text{det}} \times [K] \times [n].$$

Note that conditioning on the first relation in Theorem 15, we always have $\bar{V}_i^{k,\mu} \geq V_i^\mu$ for all $(\mu, k, i) \in \Pi^{\text{det}} \times [K] \times [n]$ because by definition $V_i^\mu = V_i^\mu(\theta^*)$.

Nash equilibrium When we choose EQUILIBRIUM in Subroutine 1 to be Nash equilibrium, by the definition of Nash-regret,

$$\begin{aligned} \text{Regret}_{\text{Nash}}(K) &= \sum_k \max_i \left(\max_{\mu_i \in \Pi_i^{\text{det}}} V_i^{\mu_i \times \pi_{-i}^k} - V_i^{\pi^k} \right) \\ &= \sum_k \max_i \left(\max_{\mu_i \in \Pi_i^{\text{det}}} \mathbb{E}_{\mu_{-i} \sim \pi_{-i}^k} \left[V_i^{\mu_i \times \mu_{-i}} \right] - V_i^{\pi^k} \right) \\ &\leq \sum_k \max_i \left(\max_{\mu_i \in \Pi_i^{\text{det}}} \mathbb{E}_{\mu_{-i} \sim \pi_{-i}^k} \left[\bar{V}_i^{k,\mu_i \times \mu_{-i}} \right] - V_i^{\pi^k} \right) \\ &= \sum_k \max_i \left(\mathbb{E}_{\mu \sim \pi^k} \left[\bar{V}_i^{k,\mu} \right] - \mathbb{E}_{\mu \sim \pi^k} \left[V_i^\mu \right] \right), \end{aligned} \tag{4}$$

where the final equality uses the fact that π^k is a Nash equilibrium of the normal-form game defined by $(\bar{V}_1^k, \dots, \bar{V}_n^k)$ as described in Subroutine 1. By Jensen's inequality and Azuma-Hoeffding inequality,

$$\begin{aligned}
& \sum_k \max_i \left(\mathbb{E}_{\mu \sim \pi^k} [\bar{V}_i^{k, \mu}] - \mathbb{E}_{\mu \sim \pi^k} [V_i^\mu] \right) \\
& \leq \sum_k \mathbb{E}_{\mu \sim \pi^k} \left[\max_i \left(\bar{V}_i^{k, \mu} - V_i^\mu \right) \right] \\
& \leq \sum_k \max_i \left(\bar{V}_i^{k, \mu^k} - V_i^{\mu^k} \right) + \tilde{O}(H\sqrt{K}) \\
& = \sum_k \max_i \left(\max_{\hat{\theta} \in \mathcal{B}^k} V_i^{\mu^k}(\hat{\theta}) - V_i^{\mu^k} \right) + \tilde{O}(H\sqrt{K}) \\
& \leq H \sum_k \max_{\hat{\theta} \in \mathcal{B}^k} \left\| \mathbb{P}_{\hat{\theta}}^{\mu^k} - \mathbb{P}_{\theta^*}^{\mu^k} \right\|_1 + \tilde{O}(H\sqrt{K}),
\end{aligned}$$

where the last equality uses the definition of \bar{V}^k and the last inequality uses the fact that the reward is an H -bounded function of the trajectory. Finally, we complete the proof by using the second relation in Theorem 15, which upper bounds $\sum_k \max_{\hat{\theta} \in \mathcal{B}^k} \left\| \mathbb{P}_{\hat{\theta}}^{\mu^k} - \mathbb{P}_{\theta^*}^{\mu^k} \right\|_1$ by $\tilde{O} \left(\frac{S^2 A O}{\alpha^2} \sqrt{K(S^2 A + S O)} \times \text{poly}(H) \right)$.

Coarse correlated equilibrium When we choose EQUILIBRIUM in Subroutine 1 to be CCE, the proof is exactly the same as for Nash equilibrium, except that the last equality in Equation (4) becomes “no larger than” by the definition of CCE.

Correlated equilibrium When we choose EQUILIBRIUM in Subroutine 1 to be CE, by the definition of CE-regret,

$$\begin{aligned}
\text{Regret}_{\text{CE}}(K) &= \sum_k \max_i \left(\max_{\phi_i} V_i^{(\phi_i \diamond \pi_i^k) \odot \pi_{-i}^k} - V_i^{\pi^k} \right) \\
&= \sum_k \max_i \left(\max_{\phi_i} \mathbb{E}_{\mu \sim \pi^k} \left[V_i^{(\phi_i \diamond \mu_i) \times \mu_{-i}} \right] - V_i^{\pi^k} \right) \\
&\leq \sum_k \max_i \left(\max_{\phi_i} \mathbb{E}_{\mu \sim \pi^k} \left[\bar{V}_i^{k, (\phi_i \diamond \mu_i) \times \mu_{-i}} \right] - V_i^{\pi^k} \right) \\
&= \sum_k \max_i \left(\mathbb{E}_{\mu \sim \pi^k} [\bar{V}_i^{k, \mu}] - \mathbb{E}_{\mu \sim \pi^k} [V_i^\mu] \right),
\end{aligned} \tag{5}$$

where the second equality uses the definition of strategy modification, and the final equality uses the fact that π^k is a CE of the normal-form game defined by $(\bar{V}_1^k, \dots, \bar{V}_n^k)$ as described in Subroutine 1. The remaining steps are the same as of the proof for Nash-regret.

B.2 Proof of Theorem 9

In this section, we prove Theorem 9 with a specific polynomial dependency as stated in the following theorem.

Algorithm 5 multi-step OMLE-Equilibrium

- 1: **Initialize:** $\mathcal{B}^1 = \{\hat{\theta} \in \Theta : \min_h \sigma_S(\hat{\mathbb{M}}_h) \geq \alpha\}$, $\mathcal{D} = \{\}$
 - 2: **for** $k = 1, \dots, K$ **do**
 - 3: compute $\pi^k = \text{Optimistic_Equilibrium}(\mathcal{B}^k)$ and **sample** μ^k **from** π^k
 - 4: **for** $h = 0, \dots, H - m$ **do**
 - 5: execute policy $\mu_{1:h}^k \circ \text{uniform}(\mathcal{A})$ to collect a trajectory $\tau^{k,h}$
 then add $(\mu_{1:h}^k \circ \text{uniform}(\mathcal{A}), \tau^{k,h})$ into \mathcal{D}
 - 6: **update**

$$\mathcal{B}^{k+1} = \left\{ \hat{\theta} \in \Theta : \sum_{(\pi, \tau) \in \mathcal{D}} \log \mathbb{P}_{\hat{\theta}}^{\pi}(\tau) \geq \max_{\theta' \in \Theta} \sum_{(\pi, \tau) \in \mathcal{D}} \log \mathbb{P}_{\theta'}^{\pi}(\tau) - \beta \right\} \cap \mathcal{B}^1$$
-

Theorem 16. (Total suboptimality of multi-step OMLE-Equilibrium) *Under Assumption 2, there exists an absolute constant c such that for any $\delta \in (0, 1]$ and $K \in \mathbb{N}$, Algorithm 2 with*

$$\beta = c \left(H(S^2 A + SO) \log(SAOHK) + \log(K/\delta) \right)$$

and EQUILIBRIUM being one of $\{\text{Nash}, \text{CCE}, \text{CE}\}$ satisfies (respectively) that with probability at least $1 - \delta$,

$$\text{Regret}_{\{\text{Nash}, \text{CCE}, \text{CE}\}}(k) \leq \tilde{\mathcal{O}} \left(\frac{S^2 A^{3m-2}}{\alpha^2} \sqrt{k(S^2 A + SO)} \times \text{poly}(H) \right) \quad \text{for all } k \in [K],$$

where the regret is computed for policy π^1, \dots, π^k .

The proof of Theorem 16 follows basically the same arguments as in the undercomplete setting, except that we replace Algorithm 4 with Algorithm 5⁴ and Theorem 15 with Theorem 17 in the first two steps. And the third step is exactly the same. To avoid noninformative repetitive arguments, here we only state Algorithm 5 and Theorem 17, while one can directly verify all the proofs in Section B.1 still hold after we make the aforementioned replacements.

Theorem 17. (Liu et al. (2022a)) *Under Assumption 2 and the same choice of β as in Theorem 9, with probability at least $1 - \delta$, Algorithm 5 satisfies that for **all** $k \in [K]$ and **all** $\theta^1 \in \mathcal{B}^1, \dots, \theta^K \in \mathcal{B}^K$*

- $\theta^* \in \mathcal{B}^k$,
- $\sum_{t=1}^k \|\mathbb{P}_{\theta^t}^{\mu^t} - \mathbb{P}_{\theta^*}^{\mu^t}\|_1 \leq \tilde{\mathcal{O}} \left(\frac{S^2 A^{3m-2}}{\alpha^2} \sqrt{k(S^2 A + SO)} \times \text{poly}(H) \right).$

We remark that Theorem 17 follows directly from instantiating the analysis of multi-step optimistic MLE (Appendix F in Liu et al. (2022a)) on Algorithm 5.

C Proofs for Playing against Adversarial Opponents

C.1 Proof of Theorem 12

In this section, we prove Theorem 12 with a specific polynomial dependency as stated in the following theorem.

⁴We remark that in each episode k of Algorithm 2, we sample the random seed ω used in π^k *only once* and then combine $\pi^k(\omega, \cdot)$ with random actions starting from different in-episode steps to collect multiple trajectories (Line 4-5 in Algorithm 2). Therefore, Algorithm 5 and Algorithm 2 are *equivalent* for the same reasons as explained in Section B.1.1.

Theorem 18. (Regret of OMLE-Adversary) *Under Assumption 1, there exists an absolute constant c such that for any $\delta \in (0, 1]$ and $K \in \mathbb{N}$, Algorithm 3 with $\beta = c(H(S^2A + SO) \log(SAOHK) + \log(K/\delta))$ satisfies that with probability at least $1 - \delta$,*

$$\sum_{t=1}^k \left(\max_{\hat{\pi}_1} \min_{\hat{\pi}_2} V_1^{\hat{\pi}_1 \times \hat{\pi}_2} - V_1^{\pi^t} \right) \leq \tilde{O} \left(\frac{S^2AO}{\alpha^2} \sqrt{k(S^2A + SO)} \times \text{poly}(H) \right) \quad \text{for all } k \in [K].$$

To begin with, for the same reasons as explained in Section B.1.2, we can directly instantiate the guarantees for optimistic MLE (Appendix E in (Liu et al., 2022a)) on Algorithm 3 and obtain:

Theorem 19. (Liu et al. (2022a)) *Under Assumption 1 and the same choice of β as in Theorem 12, with probability at least $1 - \delta$, Algorithm 3 satisfies that for **all** $k \in [K]$ and **all** $\theta^1 \in \mathcal{B}^1, \dots, \theta^K \in \mathcal{B}^K$*

- $\theta^* \in \mathcal{B}^k$,
- $\sum_{t=1}^k \|\mathbb{P}_{\theta^t}^{\pi^t} - \mathbb{P}_{\theta^*}^{\pi^t}\|_1 \leq \tilde{O} \left(\frac{S^2AO}{\alpha^2} \sqrt{k(S^2A + SO)} \times \text{poly}(H) \right).$

Now conditioning on the two relations in Theorem 19 being true, we have

$$\begin{aligned} & \sum_k \left(\max_{\hat{\pi}_1} \min_{\hat{\pi}_{-1}} V_1^{\hat{\pi}_1 \times \hat{\pi}_{-1}} - V_1^{\pi_1^k \times \pi_{-1}^k} \right) \\ &= \sum_k \left(\max_{\hat{\pi}_1} \min_{\hat{\pi}_{-1}} V_1^{\hat{\pi}_1 \times \hat{\pi}_{-1}}(\theta^*) - \max_{\hat{\theta} \in \mathcal{B}^k} \max_{\hat{\pi}_1} \min_{\hat{\pi}_{-1}} V_1^{\hat{\pi}_1 \times \hat{\pi}_{-1}}(\hat{\theta}) \right) \\ & \quad + \sum_k \left(\max_{\hat{\theta} \in \mathcal{B}^k} \max_{\hat{\pi}_1} \min_{\hat{\pi}_{-1}} V_1^{\hat{\pi}_1 \times \hat{\pi}_{-1}}(\hat{\theta}) - V_1^{\pi_1^k \times \pi_{-1}^k}(\theta^*) \right) \\ \theta^* \in \mathcal{B}^k & \leq \sum_k \left(\max_{\hat{\theta} \in \mathcal{B}^k} \max_{\hat{\pi}_1} \min_{\hat{\pi}_{-1}} V_1^{\hat{\pi}_1 \times \hat{\pi}_{-1}}(\hat{\theta}) - V_1^{\pi_1^k \times \pi_{-1}^k}(\theta^*) \right) \\ \text{by the definition of } \pi_1^k &= \sum_k \left(\max_{\hat{\theta} \in \mathcal{B}^k} \min_{\hat{\pi}_{-1}} V_1^{\pi_1^k \times \hat{\pi}_{-1}}(\hat{\theta}) - V_1^{\pi_1^k \times \pi_{-1}^k}(\theta^*) \right) \\ & \leq \sum_k \left(\max_{\hat{\theta} \in \mathcal{B}^k} V_1^{\pi_1^k \times \pi_{-1}^k}(\hat{\theta}) - V_1^{\pi_1^k \times \pi_{-1}^k}(\theta^*) \right) \\ \text{reward per episode} \in [0, H] & \leq H \sum_k \max_{\hat{\theta} \in \mathcal{B}^k} \left\| \mathbb{P}_{\hat{\theta}}^{\pi^k} - \mathbb{P}_{\theta^*}^{\pi^k} \right\|_1 \\ \text{Theorem 19} & \leq \tilde{O} \left(\frac{S^2AO}{\alpha^2} \sqrt{k(S^2A + SO)} \times \text{poly}(H) \right). \end{aligned}$$

C.2 Proof of Theorem 11




Figure 1: hard instance for Theorem 11.

The hard instance is best illustrated by Figure 1 where we sketch the transition dynamics of the POMG. Below we elaborate the construction based on Figure 1.

- **States and actions.** Each circle and rectangle in Figure 1 represents a state. Each player has two actions denoted by $\mathcal{A} = \{a_0, a_1\}$ and $\mathcal{B} = \{b_0, b_1\}$ respectively.
- **Observations.** The max-player can always directly observe the current latent state. The min-player observes the same dummy observation o_{null} in any black circle while directly observes the current state in any red circle and any rectangle. It is easy to verify by definition $\min_h \sigma_{\min}(\mathbb{O}_h) = 1$ because the emission structure is a bijection between the joint observation space and the latent state space.
- **Reward.** Only the upper rectangle emits an observation containing reward 1 for the max-player (thus reward -1 for the min-player). All other states emit observations with zero reward.
- **Transitions.** At the beginning of each episode, the environment starts from u_1 or l_1 uniformly at random. The transition dynamics from step 1 to step H only depend on the actions of the max-player while in the final step ($H \rightarrow H + 1$) the transitions are determined by the min-player. Formally,
 - When the environment is in the **upper half** of the POMG: for each step $h \in [H - 1]$, the environment will transition to $u_{h+1,0}$ if the max-player takes action a_0 and transition to $u_{h+1,1}$ if the max-player takes action a_1 . At step H , the agent will transition to the upper rectangle if the min-player takes action b_1 and transition to the lower one if the min-player picks b_0 .
 - When the environment is in the **lower half** of the POMG: for each step $h \in [H - 1]$, the environment will transition to $l_{h+1,0}$ if the max-player takes action a_0 and transition to $l_{h+1,1}$ if the max-player takes action a_1 . At step H , the agent will transition to the upper rectangle if the min-player takes action b_0 and transition to the lower one if the min-player picks b_1 .

Min-player’s optimal strategy. It is direct to see the min-player’s optimal strategy is to take action b_0 in the upper half of the POMG and action b_1 in the lower half, at step H . This strategy will lead to zero-reward for the max-player. However, implementing this strategy requires the min-player to infer which half the environment is in from her observations, which is possible only when the environment has visited some red circles in the first H steps. This is because the min-player directly observes the current state in red circles while observes the same observation o_{null} in all black circles.

Max-player’s optimal strategy. To prevent the min-player from discovering which half the environment currently lies in, the max-player’s optimal strategy is to avoid visiting any red circles.

Hardness. However, hardness happens if (a) the max-player cannot access the observations of the min-player and (b) for each $h \in \{2, \dots, H\}$, we uniformly at random pick one of $\{u_{h0}, u_{h1}\}$ and one of $\{l_{h0}, l_{h1}\}$ to be red circles, and set the remaining ones to be black. From the perspective of the max-player, she cannot directly tell which state is red or black because (a) the difference between black circles and red circles only appear in the min-player’s observations, and (b) the max-player cannot see what the min-player observes. As a result, the only useful information for the max-player to figure out which circles are red is the action picked by the min-player in the final step.

Now suppose the min-player will play the optimal strategy when she knows which half the environment is in, and pick action b_0 when she does not. In this case, for the max-player, identifying all the red circles is as hard as learning a bandit with $\Omega(2^H)$ arms where only one arm has reward $1/2$ and all other arms has reward 0. Therefore, by using standard lower bound arguments for bandits, we can show the max-player’s cumulative rewards in the first $K = \Theta(2^H)$ episodes is 0 with constant probability. In comparison, the optimal strategy, which avoids visiting all red circles, can collect $K/3$ rewards with high probability. As a result, we obtain the desired $\Omega(\min\{2^H, K\})$ regret lower bound for competing against the Nash value.

C.3 Playing against adversary in multi-step weakly-revealing POMGs is hard

In this section, we prove that competing with the max-min value is statistically hard even if (i) the POMG is two-player zero-sum and satisfies Assumption 2 with $m = 2$ and $\alpha = 1$, (ii) the opponent keeps playing a fixed action, and (iii) the player can directly observe the opponents’ actions and observations.

Theorem 20. *Assume the player can directly observe the opponents’ actions and observations. For any $L, k \in \mathbb{N}^+$, there exist (i) a two-player zero-sum POMG of size $S, A, O, H = \mathcal{O}(L)$ and satisfying Assumption 2 with $m = 2$ and $\alpha = 1$, and (ii) an opponent who keeps playing a fixed action \hat{a}_2 , so that with probability at least $1/2$*

$$\sum_{t=1}^k \left(\max_{\tilde{\pi}_1} \min_{\tilde{\pi}_2} V_1^{\tilde{\pi}_1 \times \tilde{\pi}_2} - V_1^{\pi_1^t \times \hat{a}_2} \right) \geq \Omega(\min\{2^L, k\}),$$

where π_1^t is the policy played by the learner in the t^{th} episode.

Proof. The hard instance is constructed as following:

- **States and actions:** There are four states: p_0, p_1 and q_0, q_1 . Each player has two actions, denoted by $\{a_0, a_1\}$ and $\{b_0, b_1\}$ respectively.
- **Emission and reward:** There are three different observations: o_{dummy}, o_1 and o_0 . At step $h \in [H-1]$, p_1 and p_0 emit the same observation o_{dummy} . At step H , p_1 emits o_1 while p_0 emits o_0 . Regardless

of h , q_0 always emits o_0 and q_1 always emits o_1 . Importantly, all players share the same observation. The reward function is defined so that $r(o_{\text{dummy}}) = r(o_0) = 0$ and $r(o_1) = 1$ for the max-player. Since the game is zero-sum, the reward function for the min-player is simply $-r(\cdot)$.

- **Transition:** Let x_1, \dots, x_h be a binary sequence sampled independently and uniformly at random from standard Bernoulli distribution. At step $h = 1$, the POMG always starts from state p_1 . For each step $h \in [H - 1]$:
 - If the current state is p_i , then the environment will transition to p_1 if and only if $i = 1$, the max-player plays action a_{x_h} , and the min-player plays b_0 . Otherwise, if the min-player plays b_1 , the environment will transition to q_i . Otherwise, the environment will transition to p_0 .
 - If the current state is q_i , the next state will be q_1 regardless of players' actions.

We have the following observations:

- If the min-player keeps playing b_0 , then from the perspective of the max-player the POMG essentially reduces to a multi-arm bandit problem with 2^{H-1} arms because in this case the only useful feedback for the max-player is the reward (observation) information at step H .
- The max-min (Nash) value is equal to 1, which is attained when the max-player picks a_{x_h} at step h with probability 1.
- The 2-step emission-action matrix at each step $h \in [H - 1]$ is rank 4 and has minimum singular value no smaller than 1, because we can always exactly identify the current state (for step $h \in [H - 1]$) by the current-step observation and the next-step observation if the min-player picks action b_1 in the current step.

Based on the first two observations above, we immediately obtain a $\Theta(\min\{2^H, k\})$ lower bound for competing with the max-min (Nash) value. Using the third observation, we know the POMG is 2-step weakly revealing with $\alpha = 1$, which completes the proof. \square